



Counter-hypothesis check: does focusing the model on fewer tasks lift the kept tasks above the full-baseline?

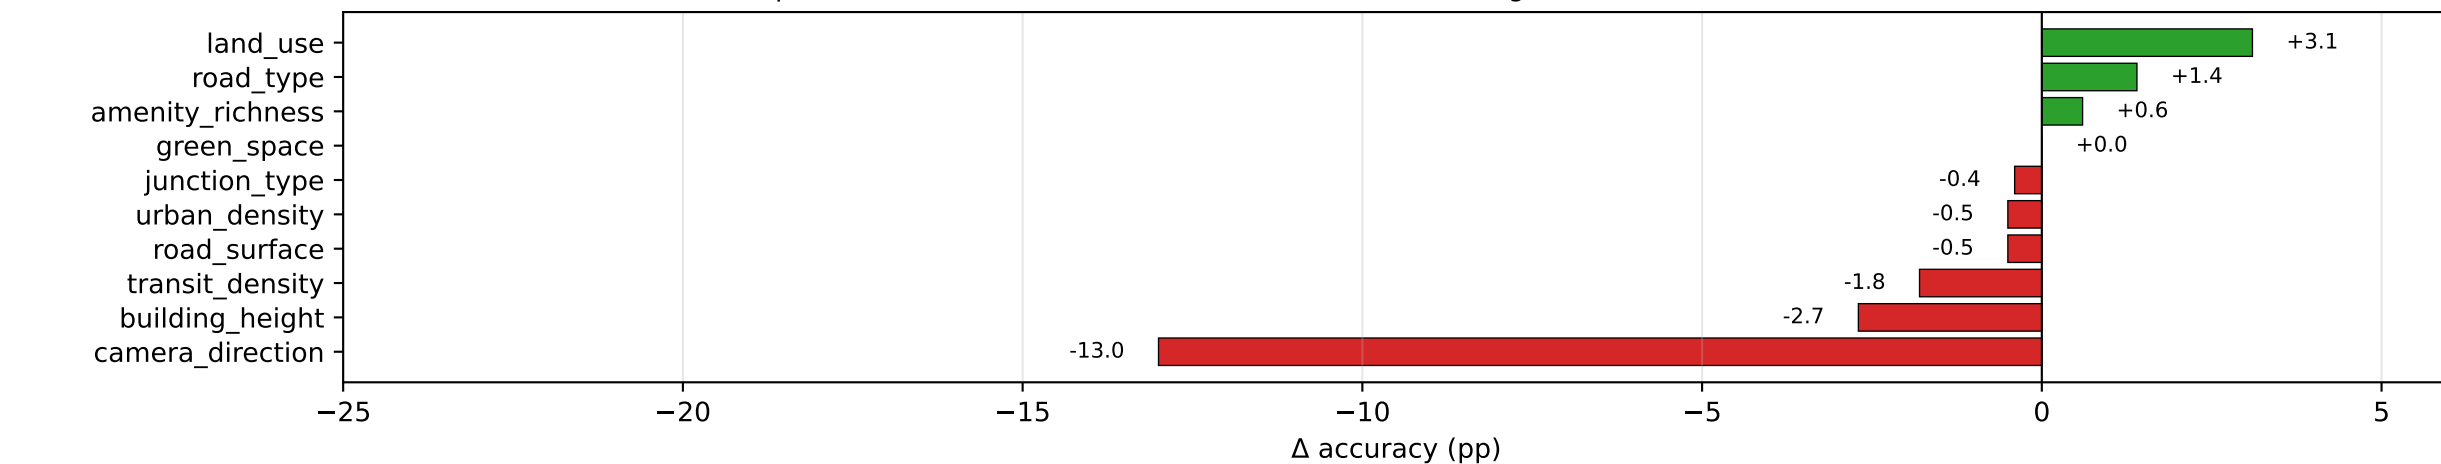
Split-1 (no urban): trained-task delta vs PC-base (negative = ablation hurt the trained task)



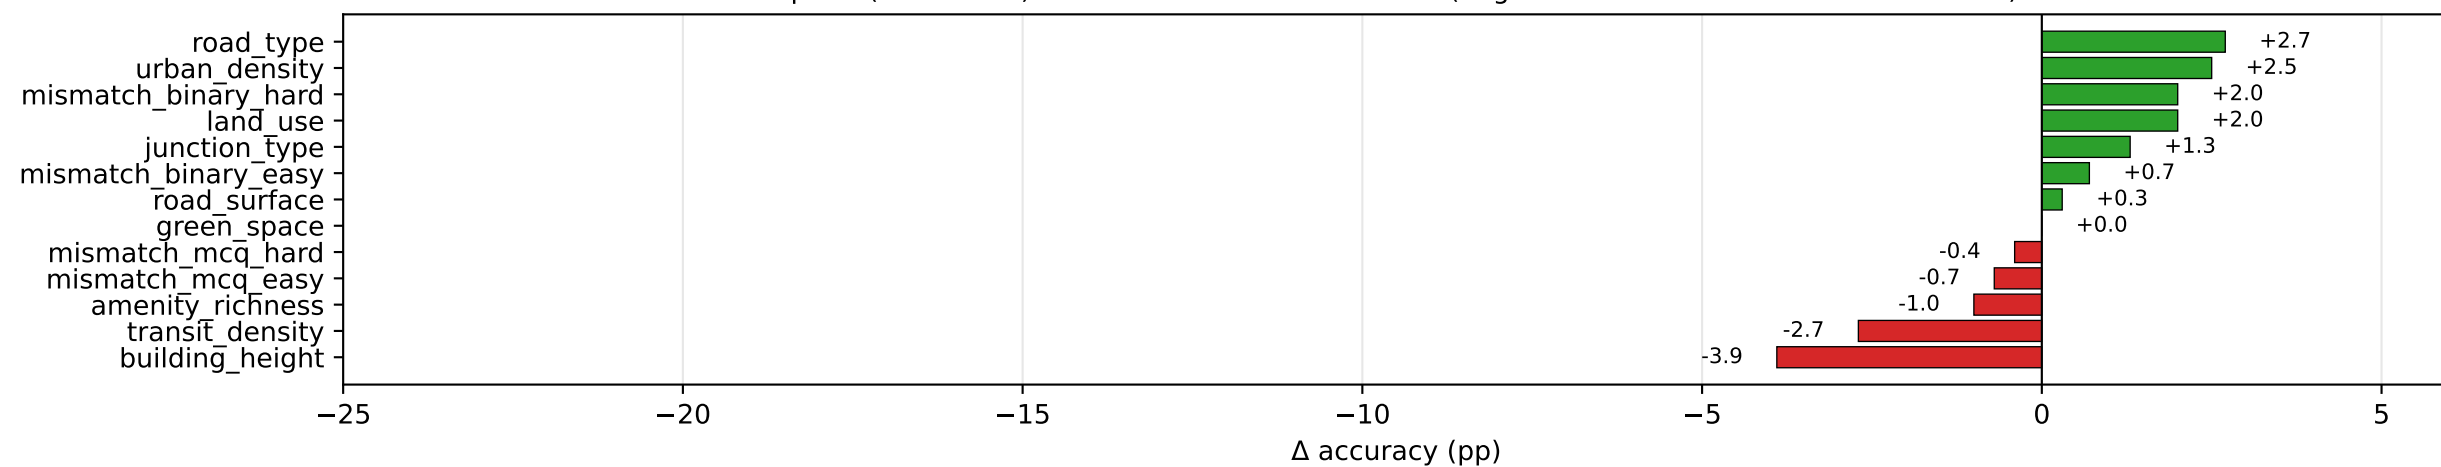
Split-2a (no geo, PC): trained-task delta vs PC-base (negative = ablation hurt the trained task)



Split-3 (no mismatch): trained-task delta vs PC-base (negative = ablation hurt the trained task)



Split-5 (no camera): trained-task delta vs PC-base (negative = ablation hurt the trained task)



Split-6 (no hard mismatch): trained-task delta vs PC-base (negative = ablation hurt the trained task)

