

Even after fine-tuning, the second view barely helps

Raw vs. SFT Qwen3.5-9B · per task · all four image modes

One table · raw vs. trained Qwen3.5-9B · 9 urbanization tasks

task	Raw 9B (zero-shot, direct)					Trained 9B (SFT, this work)				
	full	sat	sv	blind	f−b1	full	sat	sv	blind	f−b1
Land use	53.0	50.1	52.7	43.5	+0.3	72.4	71.7	69.6	67.0	+0.7
Building height	61.1	47.9	64.7	45.8	-3.6	64.7	56.3	60.5	35.3	+4.2
Urban density	38.5	39.2	44.9	40.4	-6.4	53.2	53.4	52.7	35.2	-0.2
Junction type	62.6	48.8	62.6	48.5	+0.0	76.3	71.3	72.2	51.5	+4.1
Green space	36.5	35.5	37.9	64.5	-1.4	99.0	99.5	99.5	100.0	-0.5
Amenity richness	36.8	36.1	33.5	19.5	+0.7	53.0	52.3	55.1	25.4	-2.1
Road type	47.5	40.1	57.5	34.9	-10.0	77.0	71.0	76.7	67.5	+0.3
Road surface	82.2	78.9	87.1	67.7	-4.9	95.7	95.7	94.7	95.7	+0.0
Transit density	32.1	30.9	30.4	27.3	+1.2	43.0	40.1	44.7	25.4	-1.7
OVERALL[†]	48.6	44.4	50.9	40.9	-2.3	62.5	61.1	62.0	45.2	+0.5

f−b1 = full − max(sat, sv): the gain from the *second* view. Near zero ⇒ one view already suffices.

[†]OVERALL is over the full benchmark (all topics), not just the 9 rows shown. **Fine-tuning lifts every mode by ~14 pts but does not open a fusion gap:** trained full beats the best single view by only +0.5 overall (raw: −2.3). Vision is real (full − blind = +17), but it is *one-view* vision.

What changed, what did not

- **What fine-tuning fixed:** every image mode jumps ~ 14 pts (full $48.6 \rightarrow 62.5$), and the model now clearly *uses* an image — full beats blind by **+17.3** pts (raw: +7.7).
- **What it did *not* fix:** the two-view advantage. Trained full — best-single-view = **+0.5** pt overall; on 5 of 9 tasks the best single view still ties or beats both views.
- **Conclusion:** the benchmark's questions are answerable from *either* the satellite or the street view alone. Supervised fine-tuning raises the ceiling but the second view contributes almost nothing — the same one-view story as the raw model, now at higher accuracy.